

CS745 Final Project

Choice of dataset:

Dataset is a set of surveillance videos from my ring doorbell that have been downloaded in MP4 format. The videos will be converted to image sequences for processing.



A screenshot of a Ring video feed. The video shows a man standing at a front door. A large white play button is overlaid in the center of the frame. In the top right corner, there is a "Share" button. At the bottom left, it says "0:03 / 0:25". At the bottom right, there are volume and settings icons.

	All Devices	Starred	DELETE
1	★ 🏃 MOTION - Front Door		Dec 04, 2018, 5:35 PM
2	★ 🏃 MOTION - Front Door		Dec 03, 2018, 8:12 AM
3	★ 🏃 MOTION - Front Door		Dec 02, 2018, 3:35 PM
4	★ 🏃 MOTION - Front Door		Nov 29, 2018, 1:25 PM
5	★ 🏃 MOTION - Front Door		Nov 27, 2018, 9:25 PM
6	★ 🏃 MOTION - Front Door		Nov 26, 2018, 2:05 PM
7	★ 🏃 MOTION - Front Door		Nov 23, 2018, 4:16 PM

Goal:

The goal of the project is to detect faces in the video using pre-trained classifiers from openCV and then to segment and isolate those faces within the bounding box of the detected face.

Proposed techniques:

Use of pre-trained classifiers from openCV for frontal, angled, and side view faces will be used to detect faces within the the image as well as their bounding boxes. An example of using a pre-trained openCV classifier within scikit-image can be found at: http://scikit-image.org/docs/dev/auto_examples/xx_applications/plot_face_detection.html (http://scikit-image.org/docs/dev/auto_examples/xx_applications/plot_face_detection.html). Once the bounding boxes have been identified, segmentation techniques will then be used to isolate the faces within the image and to save them to a folder of detected faces. A thresholded image will be used to isolate the detected face and covert remaining background to black.

Resources/libraries:

Two common classifiers are used for face detection in openCV LBP cascade classifier Haar feature-based classifier (uses edges/lines as features)

https://docs.opencv.org/2.4/modules/objdetect/doc/cascade_classification.html
https://docs.opencv.org/2.4/modules/objdetect/doc/cascade_classification.html,
https://docs.opencv.org/3.4/d7/d8b/tutorial_py_face_detection.html
https://docs.opencv.org/3.4/d7/d8b/tutorial_py_face_detection.html.

openCV contains a number of pre-trained models for frontal face recognition. Additional pre-trained models for side or profile faces may be downloaded and used as well.

openCV may also be used to convert video to image sequences and back to video:

https://docs.opencv.org/2.4/modules/highgui/doc/reading_and_writing_images_and_video.html
https://docs.opencv.org/2.4/modules/highgui/doc/reading_and_writing_images_and_video.html

Imports

```
In [1]: # To display inline videos in jupyter notebook
import io
import base64
from IPython.display import HTML

# For video capture
import cv2
import numpy as np

# To display real-time stills from video capture
%pylab inline
from IPython.display import clear_output

# Make full copies rather than references
from copy import copy
```

Populating the interactive namespace from numpy and matplotlib

Test Video (Indoors, Closeup with Ideal Lighting)

```
In [3]: path = "ring-videos/inside-test.mp4"
```

```
In [552]: video = io.open(path, 'r+b').read()
encoded = base64.b64encode(video)
▼ HTML(data='''<video alt="test" controls width="900">
    <source src="data:video/mp4;base64,{0}" type="video/mp4" />
</video>''.format(encoded.decode('ascii')))
```

```
Out[552]:
```

0:00

Display Video Capture in Real-Time and Save Every Keyframe

```
In [8]: ▾ def createKeyframes(path):

    # Create capture for video resource
    vid = cv2.VideoCapture(path)

    captureNumber = 0

    ▾ try:
        ▾ while(True):

            # Capture frame-by-frame
            ret, frame = vid.read()

            # Release the video resource if ret is false
            ▾ if not ret:
                vid.release()
                print("Finished saving keyframes")
                break

            # Save keyframe
            #if(captureNumber % 5 == 0):
            cv2.imwrite("keyframes/" + path.split("/")[1].split(".")[0] +

            # Convert the image from OpenCV BGR format to matplotlib RGB if
            # to display the image
            frame = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
            axis('off')
            title("Input Video, Frame %d" % (captureNumber))

            # Display the frame
            imshow(frame)
            show()

            # Display the frame until new frame is available
            clear_output(wait=True)

            captureNumber += 1

    ▾ except KeyboardInterrupt:
        vid.release()
        print("Video capture interrupted")
```

```
In [438]: createKeyframes(path)
```

```
Finished saving keyframes
```

Display Sample Keyframe

```
In [9]: ▾ def getKeyframe(file, keyframeNumber):
    keyframePath = "keyframes/" + file.split("/")[1].split(".")[0] + "_" +
    frame = cv2.imread(keyframePath)

    ▾ if frame is None:
        print("No keyframe available")
    ▾ else:
        return frame

    ▾ def displayFrame(frame, frameTitle="Keyframe"):

    ▾ try:
        convertFrame = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
    ▾ except:
        convertFrame = cv2.cvtColor(frame, cv2.COLOR_GRAY2RGB)

        axis('off')
        title(frameTitle)
        imshow(convertFrame)
        show()
```

```
In [565]: mpl.rcParams['figure.figsize'] = [8.0, 6.0]

keyframe10 = getKeyframe(path, 10)
displayFrame(keyframe10)
```



```
In [566]: keyframe85 = getKeyframe(path, 85)  
displayFrame(keyframe85)
```

Keyframe



```
In [567]: keyframe155 = getKeyframe(path, 155)  
displayFrame(keyframe155)
```

Keyframe



Experimentation with Pre-Processing

```
In [568]: checkFrame = keyframe10

displayFrame(checkFrame, "Keyframe 10, Original")

displayFrame(cv2.cvtColor(checkFrame, cv2.COLOR_BGR2GRAY), "Keyframe 10, Grayscale")

displayFrame(cv2.equalizeHist((cv2.cvtColor(checkFrame, cv2.COLOR_BGR2GRAY))), "Keyframe 10, Equalized Histogram")

clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8,8))
cl1 = clahe.apply(cv2.cvtColor(checkFrame, cv2.COLOR_BGR2GRAY))
displayFrame(cl1, "Keyframe, adaptive histogram equalization (8x8 tile)")

clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(32,32))
cl1 = clahe.apply(cv2.cvtColor(checkFrame, cv2.COLOR_BGR2GRAY))
displayFrame(cl1, "Keyframe, adaptive histogram equalization (32x32 tile)")
```

Keyframe 10, Original



Keyframe 10, Grayscale



Keyframe 10, Equalized



Keyframe, adaptive histogram equalization (8x8 tile)



Keyframe, adaptive histogram equalization (32x32 tile)




```
In [569]: checkFrame = keyframe85
num = 85

displayFrame(checkFrame, "Keyframe %d" % (num))

displayFrame(cv2.cvtColor(checkFrame, cv2.COLOR_BGR2GRAY), "Keyframe %d, Grayscale"

displayFrame(cv2.equalizeHist((cv2.cvtColor(checkFrame, cv2.COLOR_BGR2GRAY)

clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8,8))
cl1 = clahe.apply(cv2.cvtColor(checkFrame, cv2.COLOR_BGR2GRAY))
displayFrame(cl1, "Keyframe %d, adaptive histogram equalization" % (num))
```

Keyframe 85



Keyframe 85, Grayscale



Keyframe 85, Equalized



Keyframe 85, adaptive histogram equalization



```
In [570]: checkFrame = keyframe155
num = 155

displayFrame(checkFrame, "Keyframe %d" % (num))

displayFrame(cv2.cvtColor(checkFrame, cv2.COLOR_BGR2GRAY), "Keyframe %d, Grayscale"

displayFrame(cv2.equalizeHist((cv2.cvtColor(checkFrame, cv2.COLOR_BGR2GRAY)

clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8,8))
cl1 = clahe.apply(cv2.cvtColor(checkFrame, cv2.COLOR_BGR2GRAY))
displayFrame(cl1, "Keyframe %d, adaptive histogram equalization" % (num))
```

Keyframe 155



Keyframe 155, Grayscale



Keyframe 155, Equalized



Keyframe 155, adaptive histogram equalization



Detecting Faces in Keyframes

```
In [11]: face_cascade = cv2.CascadeClassifier('cascade/haarcascade_frontalface_defa  
profile_cascade = cv2.CascadeClassifier('cascade/haarcascade_profileface.x
```

```
In [33]: clahe = cv2.createCLAHE(clipLimit=2.0, tileSize=(8,8))

▼ def detectFaces(frame, show=True):
    checkFrame = clahe.apply(cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY))

    faceBoxes = []

    frontalFaces = face_cascade.detectMultiScale(checkFrame, 1.3, 5)
    #sideFaces = profile_cascade.detectMultiScale(checkFrame, 1.3, 5)

    allFaces = frontalFaces

    #if(not any(frontalFaces)):
    #    allFaces = sideFaces

    #if(any(sideFaces)):
    #    allFaces = allFaces + sideFaces

    ▼ for (x,y,w,h) in allFaces:
        faceBox = frame[y:y+h, x:x+w]
        faceBoxes.append((x,y,w,h))

        if(show):
            displayFrame(faceBox, "Face")

    ▼ if(not any(allFaces)):
        ▼ print("No faces detected")

    return faceBoxes
```

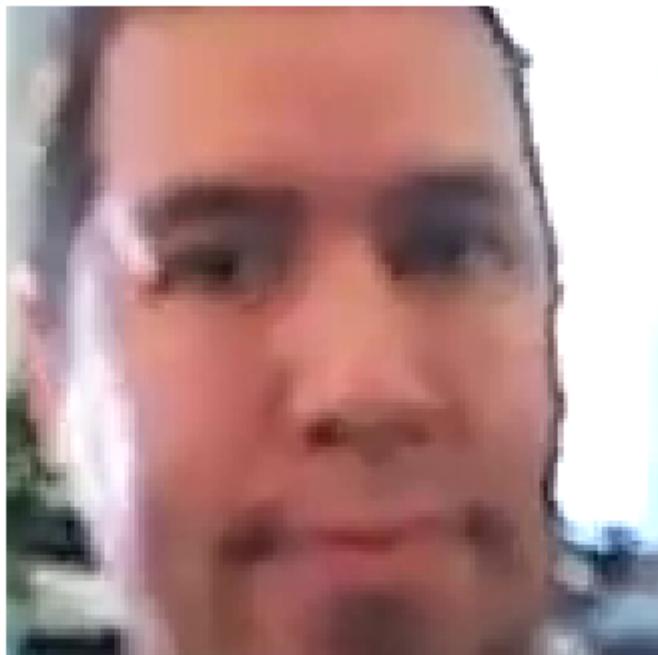
```
In [581]: detectFaces(keyframe10)
```

```
No faces detected
```

```
Out[581]: []
```

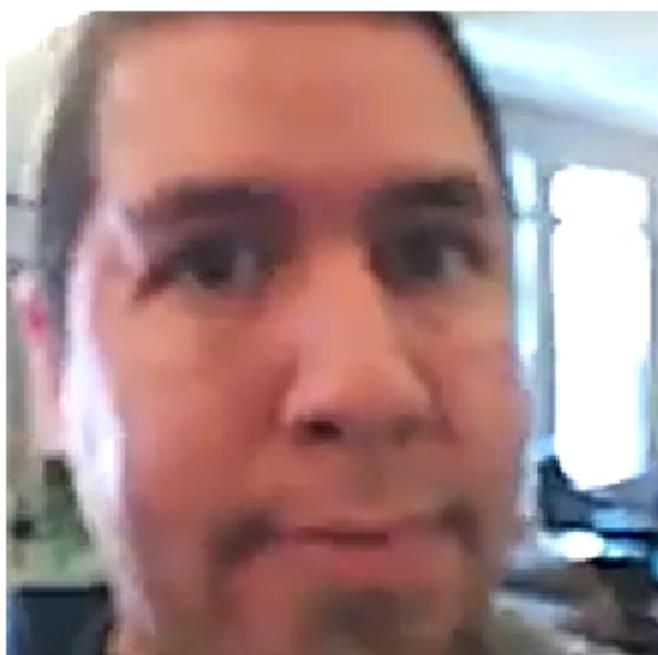
```
In [582]: _ = detectFaces(keyframe85)
```

Face



```
In [583]: _ = detectFaces(keyframe155)
```

Face



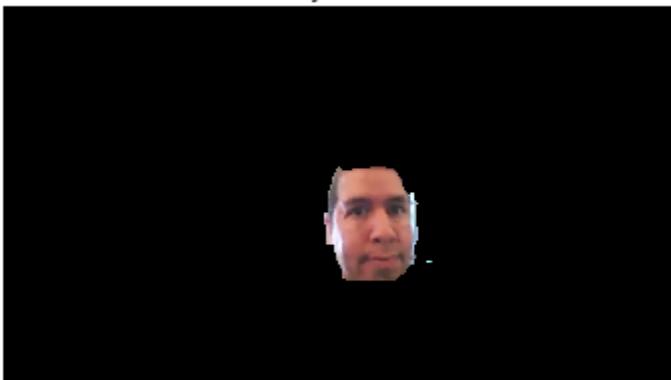
Segmenting Face

```
In [382]: ▾ for faceBox in detectFaces(keyframe155):
    mask = np.zeros(keyframe155.shape[:2],np.uint8)
    bgdModel = np.zeros((1,65),np.float64)
    fgdModel = np.zeros((1,65),np.float64)
    rect = faceBox
    cv2.grabCut(keyframe155,mask,rect,bgdModel,fgdModel,3,cv2.GC_INIT_WITH
mask2 = np.where((mask==2)|(mask==0),0,1).astype('uint8')
img = keyframe155*mask2[:, :, np.newaxis]
displayFrame(img)
```

Face



Keyframe



```
In [395]: frame = cv2.imread("two-people.jpg")
displayFrame(frame, "Example with multiple faces")

final_img = []

▼ for faceBox in detectFaces(frame):
    mask = np.zeros(frame.shape[:2],np.uint8)
    bgdModel = np.zeros((1,65),np.float64)
    fgdModel = np.zeros((1,65),np.float64)
    rect = faceBox
    cv2.grabCut(frame,mask,rect,bgdModel,fgdModel,3,cv2.GC_INIT_WITH_RECT)
    mask2 = np.where((mask==2)|(mask==0),0,1).astype('uint8')
    if(not any(final_img)):
        final_img = frame*mask2[::,:,np.newaxis]
    else:
        final_img = final_img + frame*mask2[::,:,np.newaxis]

displayFrame(final_img, "Segmented Faces")
```

Example with multiple faces



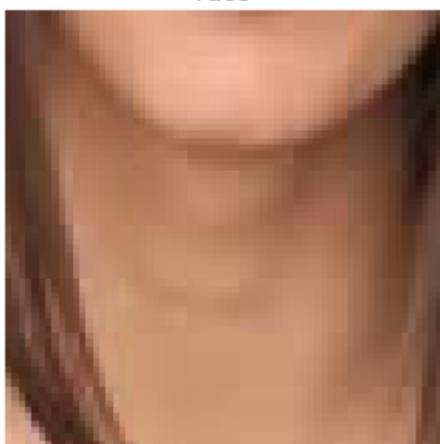
Face



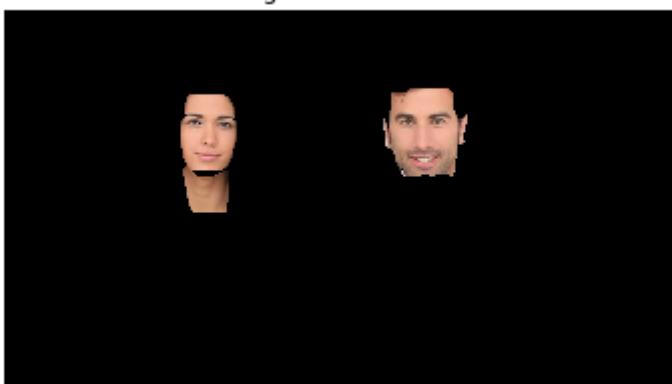
Face



Face



Segmented Faces



Creating 4 Panel Displays

1. Original
2. Preprocessed
3. Bounding box over detected faces
4. Segmented faces

```
In [12]: ▾ def display4Panel(frame, show=True):
    clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8,8))

    preprocessed = clahe.apply(cv2.cvtColor(frame,cv2.COLOR_BGR2GRAY))
    preprocessed = cv2.cvtColor(preprocessed, cv2.COLOR_GRAY2BGR)

    frameWithBox = copy(frame)
    faceBoxes = detectFaces(frame, show=False)
    ▾ for (x,y,w,h) in faceBoxes:
        frameWithBox = cv2.rectangle(frameWithBox,(x,y),(x+w,y+h),(255,0,0,2))

    segmentedFaces = []

    ▾ for faceBox in faceBoxes:
        mask = np.zeros(frame.shape[:2],np.uint8)
        bgdModel = np.zeros((1,65),np.float64)
        fgdModel = np.zeros((1,65),np.float64)
        rect = faceBox
        cv2.grabCut(frame,mask,rect,bgdModel,fgdModel,3,cv2.GC_INIT_WITH_RECT)
        mask2 = np.where((mask==2)|(mask==0),0,1).astype('uint8')
        if(not any(segmentedFaces)):
            segmentedFaces = frame*mask2[::,:,np.newaxis]
        else:
            segmentedFaces = segmentedFaces + frame*mask2[::,:,np.newaxis]

    ▾ if(not any(segmentedFaces)):
        frameWithBox = frame
        mask = cv2.cvtColor(np.zeros(frame.shape[:2],np.uint8), cv2.COLOR_GRAY2BGR)
        segmentedFaces = mask

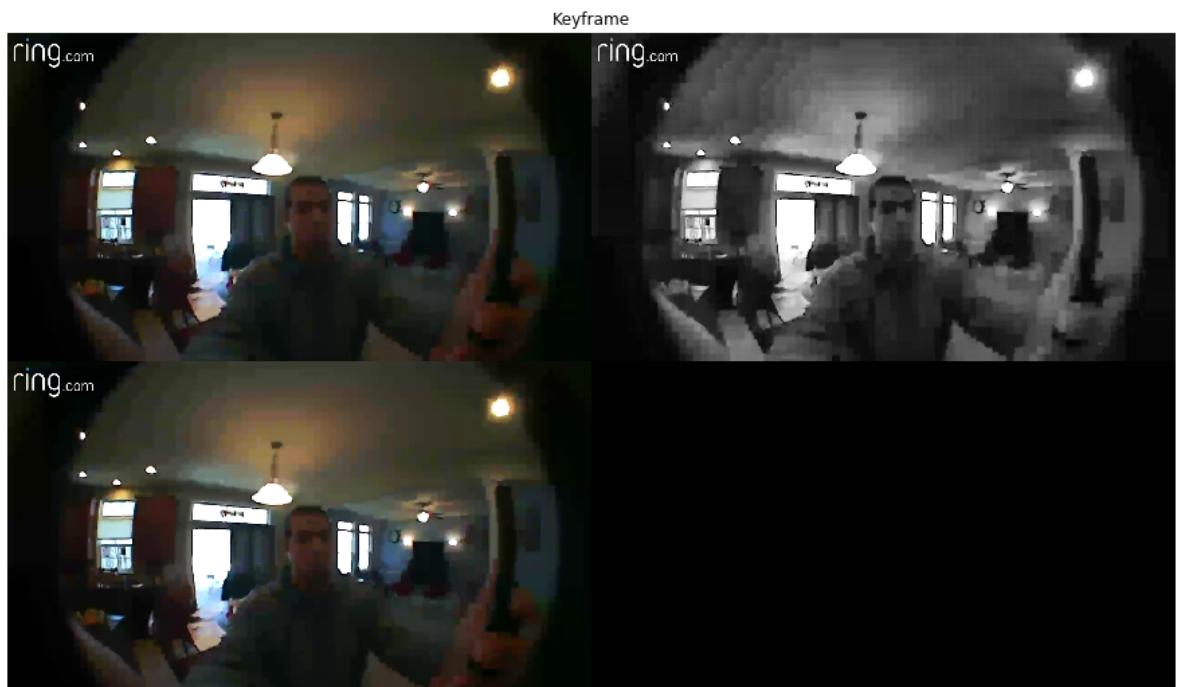
    row1 = np.hstack((frame, preprocessed))
    row2 = np.hstack((frameWithBox, segmentedFaces))
    final_frame = np.vstack((row1, row2))

    if(show):
        displayFrame(final_frame)

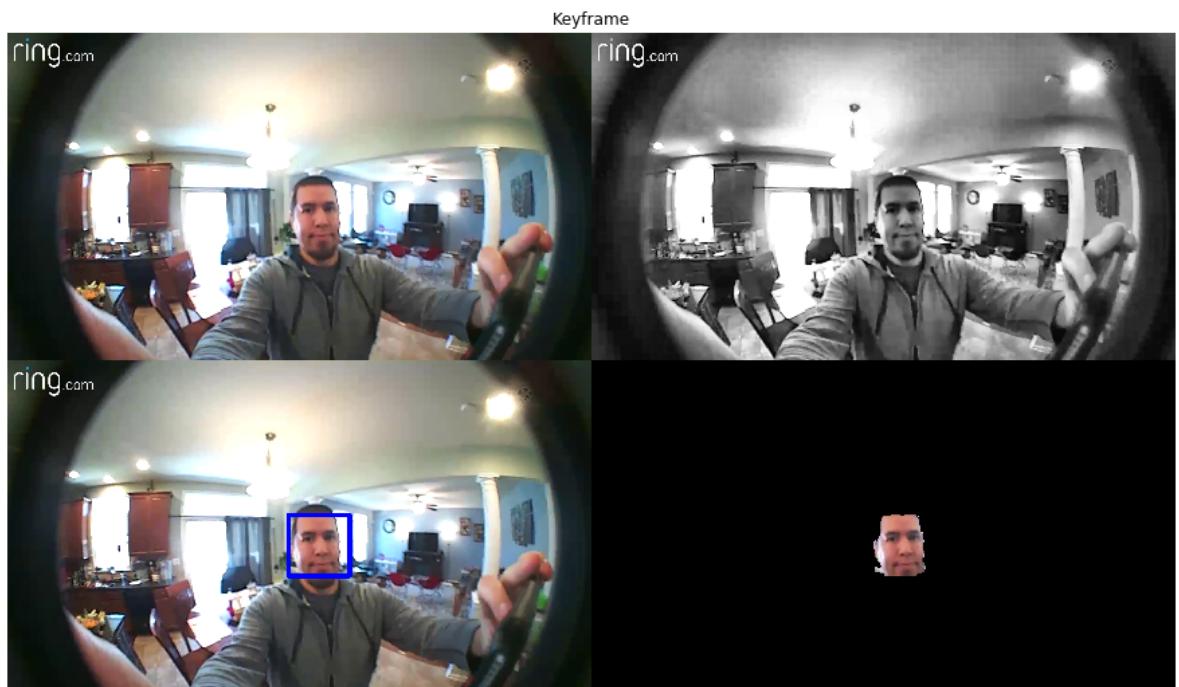
return final_frame
```

```
In [590]: mpl.rcParams['figure.figsize'] = [16.0, 9.0]
_
_=display4Panel(keyframe10)
```

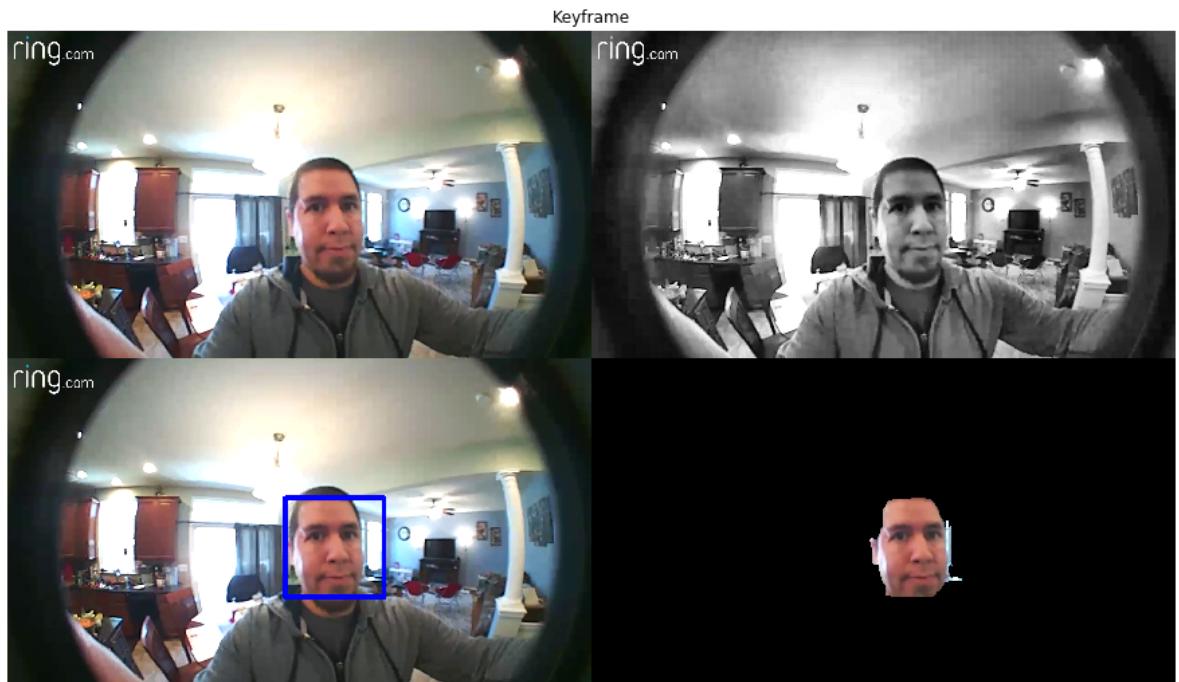
No faces detected



```
In [591]: _=display4Panel(keyframe85)
```



```
In [592]: _ = display4Panel(keyframe155)
```



Creating Output Video

```
In [17]: def createVideoFrames(path):

    createKeyframes(path)

    vid = cv2.VideoCapture(path)
    numberFrames = int(vid.get(int(cv2.CAP_PROP_FRAME_COUNT)))

    outputPath = "output-videos/" + path.split('/')[1].split('.')[0] + "_c"

    for i in range(numberFrames):
        clear_output(wait=False)
        print("Frame %d" % (i))
        frame = getKeyframe(path, i)

        panelOfFour = display4Panel(frame, show=False)
        panelOfFour = cv2.resize(panelOfFour, (0,0), fx=0.5, fy=0.5)

        cv2.imwrite("final-frames/" + path.split('/')[1].split('.')[0] + "_.jpg", panelOfFour)

    clear_output(wait=False)
    print("Finished processing image sequence, output video at %s" % (outputPath))
```

```
In [451]: createVideoFrames("ring-videos/inside-test.mp4")
          Finished saving keyframes

In [18]: createVideoFrames("ring-videos/kids-out-front.mp4")
          Finished processing image sequence, output video at output-videos/kids-out-front_output.mp4

In [21]: createVideoFrames("ring-videos/night-face-short.mp4")
          Finished processing image sequence, output video at output-videos/night-face-short_output.mp4

In [24]: createVideoFrames("ring-videos/night-face-long.mp4")
          Finished processing image sequence, output video at output-videos/night-face-long_output.mp4

In [34]: createVideoFrames("ring-videos/inside-different-angles.mp4")
          Finished processing image sequence, output video at output-videos/inside-different-angles_output.mp4

In [35]: createVideoFrames("ring-videos/inside-gray.mp4")
          Finished processing image sequence, output video at output-videos/inside-gray_output.mp4

In [36]: createVideoFrames("ring-videos/daytime-face-short.mp4")
          Finished processing image sequence, output video at output-videos/daytime-face-short_output.mp4
```

Writing Images Sequences to Video with ffmpeg

```
In [15]: ! ffmpeg -i final-frames/inside-test_%05d.png -framerate 15 -y output-video
```

ffmpeg version 4.0 Copyright (c) 2000-2018 the FFmpeg developers
built with clang version 4.0.1 (tags/RELEASE_401/final)
configuration: --prefix=/Users/Anthony/anaconda3 --cc=x86_64-apple-darw
in13.4.0-clang --disable-doc --enable-shared --enable-static --enable-zli
b --enable-pic --enable-gpl --enable-version3 --disable-nonfree --enable-
hardcoded-tables --enable-avresample --enable-libfreetype --disable-opens
sl --disable-gnutls --enable-libvpx --enable-pthreads --enable-libopus --
enable-postproc --disable-libx264
libavutil 56. 14.100 / 56. 14.100
libavcodec 58. 18.100 / 58. 18.100
libavformat 58. 12.100 / 58. 12.100
libavdevice 58. 3.100 / 58. 3.100
libavfilter 7. 16.100 / 7. 16.100
libavresample 4. 0. 0 / 4. 0. 0
libswscale 5. 1.100 / 5. 1.100
libswresample 3. 1.100 / 3. 1.100
libpostproc 55. 1.100 / 55. 1.100
Input #0, image2, from 'final-frames/inside-test_%05d.png':
 Duration: 00:00:08.40, start: 0.000000, bitrate: N/A
 Stream #0:0: Video: png, rgb24(pc), 1280x720, 25 fps, 25 tbr, 25 tbn,
25 tbc
Stream mapping:
 Stream #0:0 -> #0:0 (png (native) -> mpeg4 (native))
Press [q] to stop, [?] for help
Output #0, mp4, to 'output-videos/inside-test.mp4':
 Metadata:
 encoder : Lavf58.12.100
 Stream #0:0: Video: mpeg4 (mp4v / 0x7634706D), yuv420p, 1280x720, q=2
-31, 200 kb/s, 25 fps, 12800 tbn, 25 tbc
 Metadata:
 encoder : Lavc58.18.100 mpeg4
 Side data:
 cpb: bitrate max/min/avg: 0/0/200000 buffer size: 0 vbv_delay: -1
frame= 210 fps=184 q=31.0 Lsize= 1049kB time=00:00:08.36 bitrate=102
7.7kbits/s speed=7.33x
video:1047kB audio:0kB subtitle:0kB other streams:0kB global headers:0kB
muxing overhead: 0.166479%

```
In [20]: ! ffmpeg -i final-frames/kids-out-front_%05d.png -framerate 15 -y output-v

ffmpeg version 4.0 Copyright (c) 2000-2018 the FFmpeg developers
  built with clang version 4.0.1 (tags/RELEASE_401/final)
  configuration: --prefix=/Users/Anthony/anaconda3 --cc=x86_64-apple-darw
in13.4.0-clang --disable-doc --enable-shared --enable-static --enable-zli
b --enable-pic --enable-gpl --enable-version3 --disable-nonfree --enable-
hardcoded-tables --enable-avresample --enable-libfreetype --disable-opens
sl --disable-gnutls --enable-libvpx --enable-pthreads --enable-libopus --
enable-postproc --disable-libx264
    libavutil      56. 14.100 / 56. 14.100
    libavcodec     58. 18.100 / 58. 18.100
    libavformat    58. 12.100 / 58. 12.100
    libavdevice     58.  3.100 / 58.  3.100
    libavfilter     7. 16.100 /  7. 16.100
    libavresample   4.  0.  0 /  4.  0.  0
    libswscale      5.  1.100 /  5.  1.100
    libswresample   3.  1.100 /  3.  1.100
    libpostproc    55.  1.100 / 55.  1.100
Input #0, image2, from 'final-frames/kids-out-front_%05d.png':
  Duration: 00:00:18.72, start: 0.000000, bitrate: N/A
    Stream #0:0: Video: png, rgb24(pc), 1280x720, 25 fps, 25 tbr, 25 tbn,
  25 tbc
Stream mapping:
  Stream #0:0 -> #0:0 (png (native) -> mpeg4 (native))
Press [q] to stop, [?] for help
Output #0, mp4, to 'output-videos/kids-out-front.mp4':
  Metadata:
    encoder          : Lavf58.12.100
  Stream #0:0: Video: mpeg4 (mp4v / 0x7634706D), yuv420p, 1280x720, q=2
-31, 200 kb/s, 25 fps, 12800 tbn, 25 tbc
  Metadata:
    encoder          : Lavc58.18.100 mpeg4
  Side data:
    cpb: bitrate max/min/avg: 0/0/200000 buffer size: 0 vbv_delay: -1
frame= 468 fps=106 q=31.0 Lsize= 1499kB time=00:00:18.68 bitrate= 65
7.3kbits/s speed=4.25x
video:1496kB audio:0kB subtitle:0kB other streams:0kB global headers:0kB
muxing overhead: 0.189371%
```

```
In [28]: ! ffmpeg -i final-frames/night-face-long_%05d.png -framerate 15 -y output-  
ffmpeg version 4.0 Copyright (c) 2000-2018 the FFmpeg developers  
built with clang version 4.0.1 (tags/RELEASE_401/final)  
configuration: --prefix=/Users/Anthony/anaconda3 --cc=x86_64-apple-darw  
in13.4.0-clang --disable-doc --enable-shared --enable-static --enable-zli  
b --enable-pic --enable-gpl --enable-version3 --disable-nonfree --enable-  
hardcoded-tables --enable-avresample --enable-libfreetype --disable-opens  
sl --disable-gnutls --enable-libvpx --enable-pthreads --enable-libopus --  
enable-postproc --disable-libx264  
    libavutil      56. 14.100 / 56. 14.100  
    libavcodec     58. 18.100 / 58. 18.100  
    libavformat    58. 12.100 / 58. 12.100  
    libavdevice     58.  3.100 / 58.  3.100  
    libavfilter     7. 16.100 /  7. 16.100  
    libavresample   4.  0.  0 /  4.  0.  0  
    libswscale      5.  1.100 /  5.  1.100  
    libswresample   3.  1.100 /  3.  1.100  
    libpostproc    55.  1.100 / 55.  1.100  
Input #0, image2, from 'final-frames/night-face-long_%05d.png':  
  Duration: 00:00:18.64, start: 0.000000, bitrate: N/A  
    Stream #0:0: Video: png, rgb24(pc), 1280x720, 25 fps, 25 tbr, 25 tbn,  
25 tbc  
Stream mapping:  
  Stream #0:0 -> #0:0 (png (native) -> mpeg4 (native))  
Press [q] to stop, [?] for help  
Output #0, mp4, to 'output-videos/night-face-long.mp4':  
  Metadata:  
    encoder : Lavf58.12.100  
  Stream #0:0: Video: mpeg4 (mp4v / 0x7634706D), yuv420p, 1280x720, q=2  
-31, 200 kb/s, 25 fps, 12800 tbn, 25 tbc  
  Metadata:  
    encoder : Lavc58.18.100 mpeg4  
  Side data:  
    cpb: bitrate max/min/avg: 0/0/200000 buffer size: 0 vbv_delay: -1  
frame= 466 fps=118 q=31.0 Lsize= 1398kB time=00:00:18.60 bitrate= 61  
5.7kbits/s speed=4.72x  
video:1395kB audio:0kB subtitle:0kB other streams:0kB global headers:0kB  
muxing overhead: 0.202501%
```

```
In [27]: ! ffmpeg -i final-frames/night-face-short_%05d.png -framerate 15 -y output
```

ffmpeg version 4.0 Copyright (c) 2000-2018 the FFmpeg developers
built with clang version 4.0.1 (tags/RELEASE_401/final)
configuration: --prefix=/Users/Anthony/anaconda3 --cc=x86_64-apple-darwin13.4.0-clang --disable-doc --enable-shared --enable-static --enable-zlib --enable-pic --enable-gpl --enable-version3 --disable-nonfree --enable-hardcoded-tables --enable-avresample --enable-libfreetype --disable-openssl --disable-gnutls --enable-libvpx --enable-pthreads --enable-libopus --enable-postproc --disable-libx264
libavutil 56. 14.100 / 56. 14.100
libavcodec 58. 18.100 / 58. 18.100
libavformat 58. 12.100 / 58. 12.100
libavdevice 58. 3.100 / 58. 3.100
libavfilter 7. 16.100 / 7. 16.100
libavresample 4. 0. 0 / 4. 0. 0
libswscale 5. 1.100 / 5. 1.100
libswresample 3. 1.100 / 3. 1.100
libpostproc 55. 1.100 / 55. 1.100
Input #0, image2, from 'final-frames/night-face-short_%05d.png':
 Duration: 00:00:19.00, start: 0.000000, bitrate: N/A
 Stream #0:0: Video: png, rgb24(pc), 1280x720, 25 fps, 25 tbr, 25 tbn,
25 tbc
Stream mapping:
 Stream #0:0 -> #0:0 (png (native) -> mpeg4 (native))
Press [q] to stop, [?] for help
Output #0, mp4, to 'output-videos/night-face-short.mp4':
 Metadata:
 encoder : Lavf58.12.100
 Stream #0:0: Video: mpeg4 (mp4v / 0x7634706D), yuv420p, 1280x720, q=2
-31, 200 kb/s, 25 fps, 12800 tbn, 25 tbc
 Metadata:
 encoder : Lavc58.18.100 mpeg4
 Side data:
 cpb: bitrate max/min/avg: 0/0/200000 buffer size: 0 vbv_delay: -1
frame= 475 fps=112 q=31.0 Lsize= 1194kB time=00:00:18.96 bitrate= 51
5.8kbits/s speed=4.48x
video:1191kB audio:0kB subtitle:0kB other streams:0kB global headers:0kB
muxing overhead: 0.240515%

```
In [37]: ! ffmpeg -i final-frames/daytime-face-short_%05d.png -framerate 15 -y outrp

ffmpeg version 4.0 Copyright (c) 2000-2018 the FFmpeg developers
  built with clang version 4.0.1 (tags/RELEASE_401/final)
  configuration: --prefix=/Users/Anthony/anaconda3 --cc=x86_64-apple-darw
in13.4.0-clang --disable-doc --enable-shared --enable-static --enable-zli
b --enable-pic --enable-gpl --enable-version3 --disable-nonfree --enable-
hardcoded-tables --enable-avresample --enable-libfreetype --disable-opens
sl --disable-gnutls --enable-libvpx --enable-pthreads --enable-libopus --
enable-postproc --disable-libx264
    libavutil      56. 14.100 / 56. 14.100
    libavcodec     58. 18.100 / 58. 18.100
    libavformat    58. 12.100 / 58. 12.100
    libavdevice     58.  3.100 / 58.  3.100
    libavfilter     7. 16.100 /  7. 16.100
    libavresample   4.  0.  0 /  4.  0.  0
    libswscale      5.  1.100 /  5.  1.100
    libswresample   3.  1.100 /  3.  1.100
    libpostproc    55.  1.100 / 55.  1.100
Input #0, image2, from 'final-frames/daytime-face-short_%05d.png':
  Duration: 00:00:15.28, start: 0.000000, bitrate: N/A
    Stream #0:0: Video: png, rgb24(pc), 1280x720, 25 fps, 25 tbr, 25 tbn,
  25 tbc
Stream mapping:
  Stream #0:0 -> #0:0 (png (native) -> mpeg4 (native))
Press [q] to stop, [?] for help
Output #0, mp4, to 'output-videos/daytime-face-short.mp4':
  Metadata:
    encoder          : Lavf58.12.100
  Stream #0:0: Video: mpeg4 (mp4v / 0x7634706D), yuv420p, 1280x720, q=2
-31, 200 kb/s, 25 fps, 12800 tbn, 25 tbc
  Metadata:
    encoder          : Lavc58.18.100 mpeg4
  Side data:
    cpb: bitrate max/min/avg: 0/0/200000 buffer size: 0 vbv_delay: -1
frame= 382 fps=123 q=31.0 Lsize= 1117kB time=00:00:15.24 bitrate= 60
0.3kbits/s speed= 4.9x
video:1114kB audio:0kB subtitle:0kB other streams:0kB global headers:0kB
muxing overhead: 0.221653%
```

```
In [38]: ! ffmpeg -i final-frames/inside-different-angles_%05d.png -framerate 15 -y

ffmpeg version 4.0 Copyright (c) 2000-2018 the FFmpeg developers
  built with clang version 4.0.1 (tags/RELEASE_401/final)
  configuration: --prefix=/Users/Anthony/anaconda3 --cc=x86_64-apple-darw
in3.4.0-clang --disable-doc --enable-shared --enable-static --enable-zli
b --enable-pic --enable-gpl --enable-version3 --disable-nonfree --enable-
hardcoded-tables --enable-avresample --enable-libfreetype --disable-opens
sl --disable-gnutls --enable-libvpx --enable-pthreads --enable-libopus --
enable-postproc --disable-libx264
    libavutil      56. 14.100 / 56. 14.100
    libavcodec     58. 18.100 / 58. 18.100
    libavformat    58. 12.100 / 58. 12.100
    libavdevice     58.  3.100 / 58.  3.100
    libavfilter     7. 16.100 /  7. 16.100
    libavresample   4.  0.  0 /  4.  0.  0
    libswscale      5.  1.100 /  5.  1.100
    libswresample   3.  1.100 /  3.  1.100
    libpostproc    55.  1.100 / 55.  1.100
Input #0, image2, from 'final-frames/inside-different-angles_%05d.png':
  Duration: 00:00:15.92, start: 0.000000, bitrate: N/A
    Stream #0:0: Video: png, rgb24(pc), 1280x720, 25 fps, 25 tbr, 25 tbn,
  25 tbc
Stream mapping:
  Stream #0:0 -> #0:0 (png (native) -> mpeg4 (native))
Press [q] to stop, [?] for help
Output #0, mp4, to 'output-videos/inside-different-angles.mp4':
  Metadata:
    encoder          : Lavf58.12.100
  Stream #0:0: Video: mpeg4 (mp4v / 0x7634706D), yuv420p, 1280x720, q=2
-31, 200 kb/s, 25 fps, 12800 tbn, 25 tbc
  Metadata:
    encoder          : Lavc58.18.100 mpeg4
  Side data:
    cpb: bitrate max/min/avg: 0/0/200000 buffer size: 0 vbv_delay: -1
frame= 398 fps=107 q=31.0 Lsize= 2082kB time=00:00:15.88 bitrate=107
3.9kbits/s speed=4.29x
video:2079kB audio:0kB subtitle:0kB other streams:0kB global headers:0kB
muxing overhead: 0.122913%
```

```
In [39]: ! ffmpeg -i final-frames/inside-gray_%05d.png -framerate 15 -y output-video

ffmpeg version 4.0 Copyright (c) 2000-2018 the FFmpeg developers
  built with clang version 4.0.1 (tags/RELEASE_401/final)
  configuration: --prefix=/Users/Anthony/anaconda3 --cc=x86_64-apple-darw
in13.4.0-clang --disable-doc --enable-shared --enable-static --enable-zli
b --enable-pic --enable-gpl --enable-version3 --disable-nonfree --enable-
hardcoded-tables --enable-avresample --enable-libfreetype --disable-opens
sl --disable-gnutls --enable-libvpx --enable-pthreads --enable-libopus --
enable-postproc --disable-libx264
    libavutil      56. 14.100 / 56. 14.100
    libavcodec     58. 18.100 / 58. 18.100
    libavformat    58. 12.100 / 58. 12.100
    libavdevice     58.  3.100 / 58.  3.100
    libavfilter     7. 16.100 /  7. 16.100
    libavresample   4.  0.  0 /  4.  0.  0
    libswscale      5.  1.100 /  5.  1.100
    libswresample   3.  1.100 /  3.  1.100
    libpostproc    55.  1.100 / 55.  1.100
Input #0, image2, from 'final-frames/inside-gray_%05d.png':
  Duration: 00:00:08.68, start: 0.000000, bitrate: N/A
    Stream #0:0: Video: png, rgb24(pc), 1280x720, 25 fps, 25 tbr, 25 tbn,
  25 tbc
Stream mapping:
  Stream #0:0 -> #0:0 (png (native) -> mpeg4 (native))
Press [q] to stop, [?] for help
Output #0, mp4, to 'output-videos/inside-gray.mp4':
  Metadata:
    encoder          : Lavf58.12.100
  Stream #0:0: Video: mpeg4 (mp4v / 0x7634706D), yuv420p, 1280x720, q=2
-31, 200 kb/s, 25 fps, 12800 tbn, 25 tbc
  Metadata:
    encoder          : Lavc58.18.100 mpeg4
  Side data:
    cpb: bitrate max/min/avg: 0/0/200000 buffer size: 0 vbv_delay: -1
frame= 217 fps=127 q=24.8 Lsize= 964kB time=00:00:08.64 bitrate= 91
4.1kbits/s speed=5.04x
video:962kB audio:0kB subtitle:0kB other streams:0kB global headers:0kB m
uxing overhead: 0.182757%
```

Output Videos

Inside, Straight on, Optimal Lighting

```
In [3]: HTML(data='''<iframe width="560" height="315" src="https://www.youtube.com/
```

```
Out[3]:
```

inside test



Inside, Different Angles, Optimal Lighting

```
In [4]: HTML(data='''<iframe width="560" height="315" src="https://www.youtube.com/er
```

```
Out[4]:
```

inside different angles



Inside, Night Mode, Straight On

```
In [5]: HTML(data='''<iframe width="560" height="315" src="https://www.youtube.com/
```

```
Out[5]:
```

inside gray



Outside, Daytime, Face for Short Amounts

```
In [6]: HTML(data='''<iframe width="560" height="315" src="https://www.youtube.com/
```

```
Out[6]:
```

daytime face short



Outside, Night Mode, Face for Long Amounts of Time

```
In [7]: HTML(data='''<iframe width="560" height="315" src="https://www.youtube.com/...></iframe>')
```

Out[7]:

night face long



Outside, Night Mode, Face for Short Amounts of Time

```
In [8]: HTML(data='''<iframe width="560" height="315" src="https://www.youtube.com/...>
```

Out[8]:

night face short



Outside, Daytime, Kids from Side and My Face from Front Briefly

```
In [9]: HTML(data='''<iframe width="560" height="315" src="https://www.youtube.com/
```

```
Out[9]:
```

kids out front



```
In [ ]:
```